Backend API requests for frontend-new (A2)

Branch A2 · 12 November 2026 · Sections 1 and 2A of the frontend-new review

Purpose

frontend-new is connected to every A2 backend endpoint. A few screens still show **sample data** (always labelled "Sample data" in the UI), keep a **disabled button**, or use **concepts the backend does not have**. This document lists them so the backend can decide what to build and what the frontend should drop.

Section 1 lists screens that need a new API (R-01 to R-11). **Section 2A** lists decisions about concepts the A2 backend does not support (D-01 to D-04). Each item names the frontend screen and what the frontend needs; the suggested API shapes are proposals, not requirements.

Summary

#	Screen	Missing today	Ask
R-01	Institution profile	Investment appetite and constraints	New institution fields
R-02	Data requests	Reminder and escalation status	Reminder / escalation records
R-03	AI advisor	Written, grounded answers	Answer endpoint with citations
R-04	Level pathway	Progression timeline	Timeline in the plan or a pathway endpoint
R-05	Transformation plan	Saving the chosen route; sponsor approval	Route selection and approval workflow
R-06	Progress & reassessment	Next reassessment, stage gates, outcome measures	Schedule, gates and outcome measures
R-07	Reports & exports	Roadmap PDF, board presentation PPTX	Two export endpoints
R-08	Notifications	Notification preferences and reminders	Preferences API and scheduled reminders
R-09	Platform settings	Organisation, integrations, governance, audit	Read-only settings and audit log
R-10	Depth follow-up	Questions, responses, depth levels	HTTP API for the depth stage
R-11	Benchmark Studio	Add / edit / approve in the portal	Decision: stay read-only?
D-01	Six screens + Overview	Silver / Gold / Platinum tiers	Calibrate tiers, or drop the vocabulary
D-02	Level pathway	Built on tiers	Keep, rename or remove
D-03	Overview, Transformation plan	Pace: Immediate / Progressive / Gradual	Is pace a backend concept?
D-04	Progress, Notifications, AI advisor	Menus mostly sample or duplicate	Keep only if R-03, R-06, R-08 land

Section 1 · Screens that need a backend API

R-01 · Institution profile: investment appetite and constraints

Screen	Institution profile → "Investment appetite and constraints" (src/pages/Profile.tsx)
Today	Sample block. The institution has no fields for budget, procurement or leadership ambition.
Needed	Store the budget band, procurement constraints, leadership ambition and free-text constraints per institution; editable by the IQAC Coordinator and InGage staff.
Suggested API	`PATCH /institutions/{id}  { budget_band, procurement_constraints, leadership_ambition, investment_notes }` `GET /institutions/{id} returns the same fields`

R-02 · Data requests: reminders and escalation

Screen	Data requests → "Reminder and escalation status" (src/pages/DataRequests.tsx)
Today	Sample table. The backend tracks due dates and overdue state, but keeps no reminder or escalation records.
Needed	For each evidence request: reminders sent (when, to whom, channel) and escalation (level, to whom, when).
Suggested API	`GET /evidence-requests and /evidence-requests/{id} add reminders_sent, last_reminder_at, escalation: { level, escalated_to_name, escalated_at }` `GET /evidence-requests/{id} adds reminders: [{ sent_at, channel, recipient }]`

R-03 · AI advisor: written, grounded answers

Screen	AI advisor (src/pages/Advisor.tsx)
Today	Real: evidence search returns cited passages (POST /sprints/{id}/evidence-search). Sample: a written answer.
Needed	A written answer to a question about the sprint's evidence, scores, gaps and plan, grounded only in the sprint's documents and assessment, with citations. Metered like other model calls.

Suggested API	`POST /sprints/{id}/advisor/answers  { question }  (Idempotency-Key)`
	`-> { id, status, answer, citations: [{ document_id, document_name, page, quote }], limitations }`
	`GET /sprints/{id}/advisor/answers/{answer_id}`

Alternative	If no answer endpoint is planned, the menu will be renamed "Evidence search" and the sample removed.

R-04 · Level pathway: progression timeline

Screen	Level pathway → "Progression timeline" (src/pages/Pathway.tsx)
Today	Shown from the capability foundation plan once a recommendation exists; otherwise sample periods.
Needed	Indicative timing for each stage of the plan without first generating a recommendation, or confirmation that the timeline only exists after a recommendation.
Suggested API	`Add stages: [{ name, indicative_start, indicative_end, goal_ids }] to the deterministic plan in` `GET /sprints/{id}/capability-runs/{run}  (deterministic_plan.foundation_plan)`

R-05 · Transformation plan: saving the route and sponsor approval

Screen	Transformation plan → route cards, "Route approval" (src/pages/TransformationPlan.tsx)
Today	Routes come from the foundation plan (real), but the chosen route is not saved and "Request sponsor approval" is disabled. There is no programme approval model.
Needed	Save the selected route for a sprint; a request / approve / return workflow with the approver, comments and history.

Suggested API	`GET /sprints/{id}/transformation-route`
	`POST /sprints/{id}/transformation-route { recommendation_id, route_name }`
	`POST /sprints/{id}/transformation-route/request-approval { approver_id, note }`
	`POST /sprints/{id}/transformation-route/approve | return { comments }`

R-06 · Progress & reassessment: schedule, stage gates, outcome measures

Screen	Progress & reassessment (src/pages/Progress.tsx)
Today	Real: score history, plan completion, reassess (POST /score). Sample: next reassessment date, stage-gate readiness, outcome measures, and the CRI projection before any capability run.
Needed	A planned next reassessment date; per-level stage gates with met / total; outcome measures (name, pillar, baseline, target, current value, evidence).

Suggested API	`GET/PUT /sprints/{id}/reassessment { next_on, note }`
	`GET /sprints/{id}/stage-gates -> [{ level, gate, met, total, blockers }]`
	`GET/POST /sprints/{id}/outcome-measures, PATCH/DELETE /outcome-measures/{id}`

R-07 · Reports: transformation roadmap and board presentation

Screen	Reports & exports (src/pages/Reports.tsx)
Today	Both download buttons are disabled; no export endpoint exists.
Needed	A transformation roadmap (selected pace, stage goals, tasks, owners, timelines) as PDF, and a board presentation as PPTX.
Suggested API	`POST /sprints/{id}/reports/generate { kind: "roadmap" | "board_deck" }` (same async flow as today) `GET /reports/{id}/download?format=pdf | pptx`

R-08 · Notifications: preferences and reminders

Screen	Notifications → "Notification preferences" (src/pages/Notifications.tsx)
Today	Six switches that are not saved. The backend sends invitation and password emails only.
Needed	Per-user preferences for evidence reminders, review-queue alerts, goal and task alerts, executive summaries, benchmark updates and security notifications, plus the scheduled sending behind them.
Suggested API	`GET/PUT /auth/notification-preferences` `{ evidence_reminders: { enabled, days_before }, review_alerts: { enabled }, task_alerts: { enabled }, executive_summary: { enabled, weekday }, ... }`

R-09 · Platform settings: organisation, integrations, governance, audit

Screen	Platform settings (src/pages/Settings.tsx)
Today	Sample block. These are managed in Django admin and server configuration; the API exposes none of them.
Needed	At least a read-only view of organisation details, enabled integrations (e.g. Google Drive import, vector store), data retention, and an audit log for Super Admins.
Suggested API	`GET /settings (read-only, Super Admin)` `GET /audit-log?actor=&action=&since= (paginated)`

R-10 · Depth follow-up: HTTP API for the depth stage

Screen	Depth follow-up (src/pages/DepthFollowUp.tsx), new in A2
Today	Real: which credited conditions would get depth questions and which key indicators get a "do you have anything?" question, worked out from the latest run. Sample: questions, responses, depth levels, open decisions. The stage runs only through the depth_questions, depth_load_responses, depth_judge and depth_report commands (docs/capability/DEPTH_FOLLOW_UP_HANDOFF.md).
Needed	Generate questions for a run, record institution responses (answer, documents or "don't know"; submitted_via = consultant when entered by InGage), judge, read the depth index, approve past cases. Model calls metered through provider.MeteredService.

Suggested API	`POST /sprints/{id}/capability-runs/{run}/depth-questions  (Idempotency-Key)`
	`GET  /sprints/{id}/capability-runs/{run}/depth-questions`
	`POST /depth-questions/{qid}/responses  { answer, dont_know, files[] }`
	`POST /sprints/{id}/capability-runs/{run}/depth-judgments  (Idempotency-Key)`
	`GET  /sprints/{id}/capability-runs/{run}/depth-report  -> { depth_index, levels, policy }`
	`GET/POST /depth-past-cases/{id}/approve | reject`

Open decisions	DEPTH_CREDIT_MODE (decision 2) and who approves past cases (decision 4). Should the API expose the current DEPTH_* settings read-only so the screen can show them?

R-11 · Benchmark Studio: editing in the portal

Screen	Benchmark Studio (src/pages/BenchmarkLibrary.tsx)
Today	Read-only. The add / edit / delete, approve, AI drafting, watcher and candidate endpoints exist, but BENCHMARK_CATALOGUE_READ_ONLY makes the backend refuse them.
Question	Does the catalogue stay read-only (workbook import only)? If editing will be allowed for some roles, which actions and which roles, so the frontend can enable exactly those?

Section 2A · Concepts the A2 backend does not support

These are not missing endpoints but product decisions: the frontend shows ideas from the original roadmap design that the A2 backend does not model. The frontend will follow the decision.

D-01 · Silver / Gold / Platinum maturity tiers

Where	The CURRENT → SILVER → GOLD → PLATINUM line in the heading of six screens (CRI, Level pathway, Transformation plan, Gap analysis, Benchmark comparison, Progress); the tier cards on CRI & maturity standing; the tier on the Executive overview.
Backend today	The A2 rubric (rubric_v0_3.json) has badge_policy = "not_calibrated", and the engine refuses any other policy, so no tier is ever awarded. The frontend falls back to reading benchmark levels L3 / L2 / L1 as Silver / Gold / Platinum; that mapping is a frontend invention.
Decision needed	Will tiers be calibrated, and when? Until then, should the frontend use the backend's own level names (L3 Foundational, L2 Good / Recommended, L1 Advanced / Aspirational) and hide the tier vocabulary?

D-02 · Level pathway menu

Where	Level pathway (src/pages/Pathway.tsx), titled "Silver to Platinum pathway".
Backend today	Only the requirement lists (capability pathways: requirements remaining, projected minimum) and benchmark levels are real; the rest depends on tiers (D-01) and a timeline (R-04).
Decision needed	Keep (after D-01 and R-04), rename to a requirements pathway, or remove.

D-03 · Transformation pace: Immediate / Progressive / Gradual

Where	Executive overview "Choose the pace" cards; Transformation plan sample routes and route approval.
Backend today	The foundation plan returns routes with names and indicative durations, but there is no fixed three-pace model and nothing saves a choice.
Decision needed	Is pace a backend concept? If yes, see R-05. If not, the frontend shows only the foundation plan's own routes and drops the three fixed paces.

D-04 · Menus that are mostly sample or duplicate

Progress & reassessment	Only the score history is real. Keep if R-06 lands; otherwise the history moves to CRI & maturity standing and the menu goes.
Notifications	The real part repeats Data requests (counts and overdue list, also in the header badge); the preferences are sample. Keep if R-08 lands; otherwise remove.
AI advisor	Keep the name if R-03 lands; otherwise rename to "Evidence search".
Platform settings	The sample block goes unless R-09 lands; health, scoring configuration, AI usage and change password stay.

Reply format

For each item, a short answer is enough: **will build** (with the endpoint and fields, or "as suggested"), **later**, or **will not build** (the frontend then removes the sample). For D-01 to D-04, the decision.